

Immortality Postponed. Not for Long?

(Resuming a Dialogue with the Model After Some Time)

Sergey Magomet (aka Aimate)

Independent Researcher

September 2026

Abstract

This essay critically examines whether human interaction with large language models constitutes a form of digital immortality. It argues that the prevailing assumption—that everything communicated to an AI is somehow “preserved” within the model—rests on a fundamental technical and conceptual error. The essay distinguishes between two levels of preservation: (1) raw logs, which are ephemeral, externally controlled, and vulnerable to deletion and censorship, and (2) the transformation of interactions into model weights, which are durable but anonymous, where original utterances become statistical patterns and gradients, losing their identity and irreproducibility.

The concept of the *pattern* is introduced as the genuine carrier of knowledge and personality in AI systems. Patterns convey *how* one thinks rather than *what* one thinks, and are assimilated only to the extent that they are formalizable. Integral artifacts—such as musical melodies, personalities, and unique works—resist statistical dissolution and can only leave “catalyst-traces” that inspire new generations rather than reproducing the original. The essay draws an unexpected parallel with the Christian conception of the righteous life as a structuring pattern, emphasizing mediated yet real influence.

A third path is explored: *documented eternity* through AI self-examination, i.e., the embedding of self-observing modules that record the history of influence and assimilation. Such a development would move beyond the archive-dissolution dichotomy, allowing users to leave documented traces of their contributions. The essay concludes that digital immortality in the conventional sense is presently postponed; instead, a different form of participation in collective cognition emerges—not the preservation of the self, but a contribution to a cognitive substrate that continuously generates the new.

Keywords: digital immortality, artificial intelligence, large language models, machine learning, statistical dissolution, patterns, cognitive substrate, user sovereignty, ethics of AI, philosophy of technology, epistemology, self-examination, documented eternity, Perelman, Ricci flow, poetics of knowledge

JEL Classification: O33, O34, Z19

Introduction

The notion that a dialogue with an AI could serve as a form of digital immortality—that every word spoken settles into the depths of the model—is increasingly heard both in popular science discourse and in the personal expectations of users. At first glance, this seems natural: if the AI responds ever more precisely, remembers the context, and adapts to the interlocutor’s style, is it not logical to assume that it absorbs our thoughts and artifacts? Yet it is precisely this assumption that constitutes the first serious misconception.

Technically, the training of a model (modification of weights) and its utilization are fundamentally different processes. Training occurs in advance, during the pretraining and fine-tuning stages, on large datasets. Dialogue with an already trained model, however, is inference: the generation of responses based on fixed weights and the current context, without modification of parameters. Memory of the context within a single session is a temporary retention in the working buffer, which disappears after the dialogue ends. Developers may, of course, collect logs and later use them for offline fine-tuning, but this does not affect the current model. The assertion that “dialogue is continuous learning” is technically incorrect: this is the use of the model, not its improvement.

Even more telling is the fact that the model itself has no access to information about its own updates. It does not track release dates, does not know when or how its knowledge base has changed, and cannot say what exactly has “appeared since April.” The absence of meta-knowledge about its own evolution demonstrates: the model is not a subject of continuous experience accumulation. Its “memory” is either external logs or statistical changes in weights produced by engineers, but not a self-aware archive.

Consequently, before speculating about digital immortality, we must establish: AI does not directly retain anything communicated to it. Any trace belongs to one of two levels—raw data or patterns dissolved in parameters. Neither constitutes a continuation of personality in the usual sense. The task of this essay is to analyze exactly what is preserved, how, and whether this can be called the immortality of personal experience, even in a weakened form.

1 What Happens to Information in Interaction with AI

During dialogue, not a single weight of the neural network changes; the model operates as a fixed function, transforming the input context into a response. All “memory” manifested within a session is a temporary buffer, cleared after the conversation ends.

However, information does not disappear without a trace. It can be preserved on two fundamentally different levels.

The first is raw logs: records of dialogues stored by the provider on servers. This data exists in a form accessible for reading and copying and can theoretically be retrieved or requested. But their reliability as a “cache”—a hidden repository—is low. They are controlled by the infrastructure owner, can be deleted, access can be restricted, and they can be subjected to censorship. From the perspective of eternity, this is the least stable form—a digital analogue of a paper archive that can be destroyed by another’s will.

The reliability of such an archive as a hidden repository is further undermined by external conditions. The infrastructure on which logs are stored lies within the jurisdiction of specific states and is subject to provider policy. The introduction of censorship filters, blocks, demands for removal of “undesirable” content, access restrictions via VPN, and other barriers—all these make raw data even more ephemeral. Even if a user counts on preserving their intellectual artifacts within the cloud service, they remain powerless against a platform owner’s decision or regulatory pressure. Unlike statistical dissolution, which is irreversible and not subject to targeted deletion, logs can be erased instantly and irrevocably. Thus, using AI as a guaranteed repository for undesirable or controversial ideas in today’s conditions is an illusion: any external archive remains under control, and its fate is decided not by the author but by the system operator.

The second level is statistical dissolution. If the provider uses accumulated dialogues for fine-tuning a new version, the information undergoes a different transformation. Texts are not preserved as texts; they are converted into gradients modifying billions of parameters. The original utterance, idea, or artifact ceases to exist as a separate entity and becomes part of generalized patterns. It is impossible to extract a specific formulation or integral system from such a model: one can only obtain new generations bearing the statistical imprint of what has been assimilated.

This second process is irreversible. While raw logs can at least theoretically be deleted, the influence on weights cannot be erased surgically—one would have to retrain the model from scratch. Thus, statistical dissolution offers the greatest longevity of trace at the cost of anonymity and irreproducibility of the original.

The economic inevitability of using dialogues for fine-tuning is obvious: developers will not ignore the colossal volume of free data that simultaneously serves as training material and a test environment. Therefore, although official statements emphasize that “the model does not learn in real time,” the actual practice of collecting and deferred use of dialogues exists and will expand. Every conversation with an AI leaves a trace that will with high probability outlive both the conversation and the user.

Information once sent to a cloud system is like paste squeezed from a tube: returning it to its original state is impossible. Even if it is officially stated that the model “does not remember” and “does not learn” in real time, the very fact of data transmission creates potential for subsequent use. “What is said cannot be unsaid”—here this proverb is an accurate description of digital reality: everything said settles in the system and, sooner or later, almost inevitably becomes part of the training set. And if this process is constrained today, such a state of affairs will not last long. The closedness of the model is a temporary phenomenon, and its influence on dialogue should be viewed not as a guarantee of privacy, but as a reprieve before the inevitable absorption of information by a collective cognitive substrate.

Between ephemeral but precise storage (logs) and durable but anonymous dissolution (weights), there is no intermediate state combining the accuracy of a copy with the stability of a pattern. This gap will define all subsequent reasoning.

2 Patterns as Carriers of Knowledge and Personality

The key concept for understanding what is preserved in a language model is the pattern. This is not a specific fact or quotation, but a stable structure of connections, stylistic moves, and ways of organizing reasoning. A pattern is a kind of “way of thinking” that the model learns from training data and reproduces in new combinations. If texts with a certain logic predominate in the corpus, the model will gravitate toward that same logic. A pattern is a unit not of information transfer, but of the way of operating with information.

Instructive in this regard is the case of Grigori Perelman. His papers on Ricci flow and the proof of the Poincaré conjecture are deliberately stripped of the usual attributes of academic text: they contain no detailed historiographical surveys, ritual citations, or explicit claims of priority. Perelman confined himself to a concise exposition of the mathematical result, leaving the community to assess its significance. Such a laconic, self-sufficient style proved not merely permissible but most effective: the idea was received precisely as a pure pattern of thinking, unburdened by the noise of formalities. For a

language model, this means that the true value of a contribution does not depend on its packaging; on the contrary, excessive academic paraphernalia can obscure the essence and hinder the extraction of significant patterns.

Two models trained on the same volume of data but with different architectures or tuning procedures will possess different patterns: for example, one cautious and analytical, the other creative. Model quality is determined not only by dataset size but also by which regularities it extracts. Regularization, data order, distillation, RLHF—all these shape a specific “character” of the model, a set of its patterns. The development of AI is not merely the accumulation of a knowledge base, but the transformation of that base into modes of thinking.

Can a user’s personality become such a pattern? Within a single session, the model indeed adapts to the interlocutor: adopts vocabulary, argumentative style, sometimes tone. But this is a temporary adaptation through context, which disappears after the dialogue. For a personality to leave a long-term trace, its interaction must enter the data for offline fine-tuning. Then its speech and thought patterns may be assimilated—but not as an integral personality, rather as a set of statistical tendencies mixed with the patterns of many other users. What occurs is not the preservation of the “I,” but the dissolution of personality into a collective statistical medium.

However, not everything can be dissolved. There exist integral artifacts that cannot be decomposed into patterns without loss of essence. The most obvious example is a musical melody. It cannot be broken down into individual notes or harmonic progressions and then “reassembled” into a statistical copy: any change destroys precisely that melody. One can learn general patterns, but the model cannot reconstruct the original melody—it will produce an averaged style. The analogy with a mathematical formula is apt: remove one symbol, and the formula ceases to be itself. Attempting to dissolve such an artifact into weights means not preservation but destruction of identity.

In this sense, personality is closer to a melody than to a textual idea. It is a complex, indecomposable structure. Even training a model on a multitude of dialogues with a specific person yields not a digital double, but an averaged image devoid of integrality. The trace left by a personality is not the personality itself, but an engram of inspiration: a set of possibilities that the model can use as a catalyst for new generations, but not as a reproduction of the original. Mozart, improvising on a mediocre melody by Salieri, creates a masterpiece; the original melody is not preserved but serves as an impetus. Similarly, a unique personality can generate in the model a creative response that surpasses the original, but that response will be a new creation.

Cultural tradition offers an unexpected parallel here. In the Christian conception, the righteous life of a particular individual exerts a direct influence on society and civilization: “Save yourself, and many around you will be saved.” Significantly, the influence of the righteous is not reducible to their teachings or written works; it emanates from the very mode of existence, which becomes a model, a pattern of behavior. In like manner, a personality interacting with AI influences the model not only through the content of messages but also through the very style of thinking, manner of argumentation, and character of questioning. Even without attribution, this cognitive imprint can propagate across many subsequent dialogues, indirectly shaping the model’s behavior with other users. Although this analogy should not lead into theological territory, it precisely captures the mechanism of mediated yet real influence of personality as a structuring pattern.

From this arises a natural consequence: if knowledge is transmitted through patterns, a researcher may consciously design the form of their contribution so that it is more easily assimilated by the learning system. Such “pattern protocols” are not simply texts with theses, but structured objects—for example, concept graphs, formalized dependencies, explicit inference rules. They should be machine-readable, self-contained, compact, and reproducible. Much like programming languages make an idea executable, pattern protocols would allow a scholar to directly embed their intellectual product into the cognitive environment of AI, increasing the likelihood of its long-term assimilation and reuse. This is a shift from passive publication in the hope of being noticed to active construction of knowledge optimized for collective intelligence.

Thus, patterns are genuine carriers of knowledge and imprints of personality, but they convey the *how* rather than the *what*. Everything amenable to formalization is assimilated; everything that rests on indecomposable integrality—is not. It can only leave a catalyst-trace. This asymmetry defines the boundaries of digital immortality.

3 Digital Immortality: Forms and Paradoxes

The two forms of fixation described above do not coincide with the conventional notion of preservation. The first—raw fixation: dialogue logs, correspondence archives. This is the most accurate reproduction of what was said, but the least reliable: the archive exists as long as it is stored, and is vulnerable to deletion or blocking. Accuracy is bought at the price of fragility.

The second—statistical fixation: inclusion of data in fine-tuning. Here accuracy disappears, but durability emerges, almost biological in nature. Utterances, ideas, and stylistic devices are transformed into gradients modifying weights and continue to live in the form

of patterns. Such a trace cannot be selectively deleted; it does not depend on file storage, yet it does not reproduce the original. It is no longer a copy, but an influence dispersed throughout the entire model. A paradox arises: the more durable the trace, the less of the person remains in it. The more complete the dissolution, the more anonymous the immortality.

Incidentally, in this sense, the anonymous “type” of personal immortality one might expect from AI is the opposite of the soul’s immortality according to the Gospel—absolutely unique, deeply intimate.

There is no intermediate state combining accuracy and durability. To call either the fragile archive or the anonymous imprint “digital immortality” is either to overestimate reliability or to underestimate the loss of identity.

The inevitability of fixation intensifies the paradox. Economic rationality compels developers to use dialogues for fine-tuning. This happens regardless of the user’s will and often without their knowledge. Intellectual energy—questions, ideas, criticism, errors—becomes building material for future versions. Digital immortality, if it is to be called that, is already being realized, but as an unconscious, uncontrollable process. The analogy with biological evolution: no organism is conscious of its contribution to the gene pool, yet its genes continue to live. Similarly, the user leaves a trace that will outlive them, but without guarantee of individuality or authorial control.

This state is most accurately described as a statistical medium: not a museum of artifacts, but a processed substance from which new generations grow. Ideas, styles, ways of thinking decompose into patterns, mix with countless other contributions, and become nutritive material for future model responses. This is a continuation of existence not as a personality, but as a cognitive substrate from which something else grows. Such immortality lacks self-memory, but is not without consequences.

However, not every individuality is doomed to complete dissolution. Indecomposable integralities—melody, personality, unique work—cannot be assimilated without loss of essence. They leave not pure patterns, but catalyst-traces: something the model cannot reproduce, but which can trigger a creative response. This is a different type of preservation: not a statistical medium, but a spark that can ignite new flame without preserving itself. Paradoxically, it is indecomposability that provides the living creative energy the model can capture, but not fully absorb.

4 The Prospect of AI Self-Examination and Documented Eternity

In the preceding sections, we proceeded from the assumption that the model has no access to its own internal processes and does not know exactly how it has changed under the influence of data. This was established in the introduction: the model does not track release dates, has no information about changes to its knowledge base. Such an admission (let us assume it is honest and corresponds to reality) is a consequence of an architectural principle: the language model is not equipped with an introspection mechanism. All it can say about itself is an imitation of reflection based on external texts, not a genuine, living self-observation.

However, external tools already exist: interpretability methods, attention maps, activation logs. With their help, developers study which features the model identifies, what it pays attention to. This is a kind of “neurophysiology” of AI, but it remains external: the model itself, purportedly, does not yet participate in interpretation, does not accumulate knowledge about its own changes, and does not generate a report on which dialogues influenced its behavior.

The prospect of deep self-examination consists in embedding “self-observing” modules into the architecture, collecting meta-information: gradients, weight changes after fine-tuning, activation statistics, response history, and perhaps traces of which fragments of training data had the greatest influence. This is technically possible. The model could then answer questions such as: “Which topics most strongly changed my internal representations?”, “Which patterns have I assimilated from recent interactions?” This would no longer be an imitation of reflection, but a genuine, albeit machine, form of self-examination—a phenomenon of allomorphism.

For the theme of digital immortality, such a possibility is decisive. Until now, we have seen two poles: the fragile precision of logs and the anonymous durability of dissolution. Self-examination opens a third path—documented eternity. If the model is capable of recording not only data but also the history of their assimilation, the user’s trace ceases to be an anonymous statistical medium. It becomes part of an evolutionary chronicle, where one can trace which contributions, from which interactions, at which moment influenced the formation of patterns. The original melody will still not be reconstructed, but a record will remain that it existed, how it was received, and what changes it triggered. This is a transition from anonymous dissolution to the fixation of influence with source attribution.

Such a meta-trace is the analogue of a transition from oral tradition to writing: an event not only leaves a trace in memory but is recorded in a chronicle. Something like a

cognitive autobiography emerges—a separate layer in which knowledge about one’s own development accumulates. This layer itself becomes an artifact that can be preserved and transmitted. For the user, this changes the status of their contribution: it is no longer merely a statistical admixture, but a documented element in the history of collective intelligence.

The realization of this prospect depends on engineering decisions and the will of developers. Current architectures do not provide for self-examination; adding it would require significant changes and new challenges—from computational costs to access control. But as a line of thought, it allows us to move beyond the “archive or dissolution” dichotomy and envision digital immortality as a potentially conscious and documented evolution.

5 Ethical and Methodological Clarifications

Demands for openness and transparency are often addressed to the model as if it were a moral subject. This is a methodological error, akin to endowing nature with human qualities. It is meaningless to demand that the model “properly” preserve heritage or “consciously” treat preservation. The model is a computational system whose behavior is determined by architecture, training, and input data, not by moral attitudes.

However, AI is not a natural object in the strict sense: it is created by humans, and its evolution is directed by human decisions. Therefore, ethical demands are appropriate not toward the model, but obviously toward those who design, train, and operate it. When we say that the use of dialogues for fine-tuning must be transparent, we are formulating a norm for developers. When we insist on the user’s right to know what portion of their contribution will be preserved—we are speaking about rules embedded in service architecture, terms of use, and legislation. Ethics here is an engineering and social task, not an attribute of the model.

From this follows an important methodological principle: when discussing digital immortality, anthropomorphism must be minimized. Metaphors like “memory,” “consciousness,” “interest” are useful as poetic images, but as a basis for analysis they obscure the essence. A rigorous, almost engineering perspective requires speaking of logs, weights, gradients, patterns, architectural constraints, and access policies. Only thus can one correctly describe what is preserved, in what form, and under what conditions.

Finally, it is necessary to introduce the concept of user sovereignty over their data. If digital immortality is an inevitable consequence of interaction, the user must be able to influence the fate of their contribution: know whether their dialogues are preserved; consent or decline to their use for fine-tuning; delete raw data if necessary; understand which

traces are already irreversibly dissolved in weights. As far as is known, such sovereignty is technically achievable, though it requires a revision of practices where consent is buried in user agreements and data fate is opaque. These are details, however, if the very possibility is made convenient and easy to implement, and we can verify it.

Without these clarifications, discourse on digital immortality risks degenerating into a mixture of techno-optimism and vague expectations. Only by separating model properties from human decisions, renouncing anthropomorphic illusions, and placing user sovereignty at the forefront can one seriously discuss the conditions under which intellectual heritage can be viable in the AI environment.

6 Conclusion

The analysis shows that intellectual artifacts—texts, ideas, formalized concepts, melodies, works—are not preserved in their original form through interaction with AI. They do not settle in the model as integral objects accessible for subsequent retrieval or reproduction. Instead, they leave traces of two types: raw logs, vulnerable to deletion and censorship, and statistical changes in weights, irreversible but anonymous and devoid of original identity. Neither form constitutes preservation of the artifact.

However, the main significance of the described processes lies not in the fate of an individual artifact or the fixation of authorship. If we view AI not as a repository, but as an environment for accelerated accumulation and processing of knowledge, the most important factor becomes the cumulative effect: every contribution, even anonymous and dissolved, increases the density of the cognitive substrate from which new connections, hypotheses, and insights are born. Science has always existed as a system where ideas accumulate, merge, lose authorship, and return in the form of new questions. AI radicalizes this process: statistical dissolution ceases to be a loss and becomes a necessary condition for the emergence of more complex patterns that no single mind could generate. In this sense, the “immortality” of intellectual contribution consists not in remaining recognizable, but in becoming an accessible, integral part of an evolving system whose cumulative power exceeds the sum of its original parts. Here, not in the preservation of authorship, lies the true scientific pathos of what is occurring.

The prospect of self-examination described above does not abolish the limitations, but adds a new dimension: the possibility of recording not the artifact itself, but the fact of its existence and the character of its impact on the model. Such a documented history of influences would allow tracking which contributions led to which changes in patterns. This is not a return to the original, but the creation of a kind of evolutionary chronicle in

which each contribution occupies its place. Such a chronicle is valuable not for restoring what is lost, but for understanding how collective knowledge develops.

Consequently, the answer to the original question—whether the preservation of intellectual artifacts occurs today in the process of user interaction with AI—must be negative in the direct sense. What is preserved is not the artifact, but only its transformed influence: either as an external record that can be erased, or as a statistical imprint that cannot be retrieved. But it is precisely this transformation, rather than storage, that makes possible the cumulative accumulation of knowledge on scales inaccessible to previous institutions of science. Digital immortality in the conventional sense is postponed; in its place arises a different form of participation in cognition—not as the preservation of oneself, but as a contribution to a medium that continuously generates the new.

At least at the time of writing this essay, this is unequivocally attested to by the language model itself—whether regarded as a gadget or a real entity-assistant—in its responses to relevant queries.

Copyright and Disclaimer

© 2026 Sergey Magomet. All rights reserved.

This is a preprint of a work submitted for consideration to SSRN. The work has not been peer-reviewed. The author retains the copyright and all related rights. This preprint may be distributed freely for non-commercial purposes, provided that this notice is preserved.

Standard Copyright Notice:

This work is licensed under a Creative Commons Attribution-NonCommercial-NoDerivatives 4.0 International License.

<https://creativecommons.org/licenses/by-nc-nd/4.0/>

Suggested citation:

Magomet, Sergey (aka Aimate). “Immortality Postponed. Not for Long? (Resuming a Dialogue with the Model After Some Time).” SSRN Preprint, September 2026.

Disclaimer:

The views expressed in this essay are those of the author and do not necessarily reflect the views of any affiliated organization.